

Charles Leung

New York, NY | leung.charles@outlook.com | (929) 266-6089 | linkedin.com/in/charles-leung

Executive Summary

AI Engineering Manager (ex-Founding Data Engineer) with 11 years building distributed data platforms and production machine learning systems. Built greenfield data and ML infrastructure from 0-to-1 through EvolutionIQ's \$730M acquisition by CCC Intelligent Solutions.

- **Applied AI & ML:** Multi-Agent Workflows, LLM-as-a-Judge, Eval Harnesses, Inter-Annotator Agreement, Small-Model Self-Hosting (Gemma, vLLM), Vertex AI, Retrieval Augmented Generation (RAG)
- **Data & Distributed Systems:** Dagster, Kafka, ksqldb, BigQuery, PostgreSQL, duckdb, pandas, numpy, numba, polars, Apache Airflow, Docker, Kubernetes, GCP, AWS, Temporal
- **Leadership & Delivery:** Multi-Carrier Enterprise Delivery, Cross-Functional Gating (Solutions & QA), Release Governance, Hiring Loop Design, Technical Mentorship.

EvolutionIQ (\$730M Acquisition by CCC Intelligent Solutions)

Engineering Manager — Applied AI & Data | Nov 2023 – Present

- Led the Workers' Comp engineering execution across the stack, shipping a total of 26 product modules across 6 Workers Compensation Carriers.
- Designed our deployment process from initial sales calls to go-live, using AI Agents to communicate status across JIRA via timelines and Fix Versions, leading to significantly improved predictability in our launch timelines.
- Ran small-model R&D (quantized Gemma / local models) to match commercial LLM benchmarks for claim extraction. Partnered with platform engineering to self-host inference in production, reducing tail latency and cutting token expenditure across key carrier accounts.
- Built an internal agentic manager tool (LangGraph + Vertex AI) that parses GitHub, Jira, and Slack activity into calibrated performance review packets. Cut year-end review prep from multi-day crunches to 3 hours; adopted by many peer engineering managers.
- Built evaluation harnesses and ground-truth validation pipelines for LLM document extraction and automated claim guidance. Managed annotator pools on experiment design (sampling strategies, inter-annotator agreement)
- Led the batch-to-realtime transition across the organization by authoring technical decision proposals, cleared InfoSec compliance, and worked with a partner team to deliver an end-to-end pipeline into production.

Founding Data Engineer & Tech Lead | Oct 2021 – Nov 2023

- Built greenfield data and ML infrastructure from 0-to-1 for the Workers' Compensation vertical, scaling architecture through company growth to a \$730M acquisition by CCC Intelligent Solutions.
- Redesigned batch processing infrastructure using Change Data Capture (CDC) patterns and tuned Kubernetes resource allocation on GCP, cutting end-to-end data processing latency by 96% (7 hours to 16 minutes).
- Implemented centralized feature store and ML pipeline architecture with Dagster and BigQuery, increasing feature experimentation velocity and model iteration speed by 80%.
- Designed and productionized an automated vocational recommendation engine and a tree-based crosswalk algorithm across clinical ICD code versions.
- Founded the internal Data Guild, establishing DBT, Dagster, and modular SQL design patterns

Tradeweb Markets — Data Platform Engineer (Mar 2020 – Oct 2021)

Barclays Investment Bank — Quantitative Equity Developer (Oct 2018 – Mar 2020)

Société Générale CIB — Rapid Application Developer (Jan 2017 – Oct 2018)

EVRAZ Inc. NA — Steelmaking Process Engineer (Jun 2013 – May 2016)

Carnegie Mellon University — Bachelor of Science in Chemical Engineering | Bachelor of Science in Technical Writing